

New Computer Trends And How This Affect Us

Francesc Alted

Freelance Consultant

<http://www.blosc.org/professional-services.html>



April 10rd, 2016

“No sensible decision can be made any longer without taking into account not only the world as it is, but the world as it will be.”

— *Isaac Asimov*

“No sensible decision can be made any longer without taking into account not only the **computer** as it is, but the **computer** as it will be.”

— My own rephrasing

About Me

- Physicist by training
- Computer scientist by passion
- Open Source enthusiast by philosophy
 - PyTables (2002 - 2011)
 - **Blosc** (2009 - now)
 - bcolz (2010 - now)

Why Open Source Projects?

“The art is in the execution of an idea. Not in the idea. There is not much left just from an idea.”

–Manuel Oltra, music composer

“Real artists ship”

–Seth Godin, writer


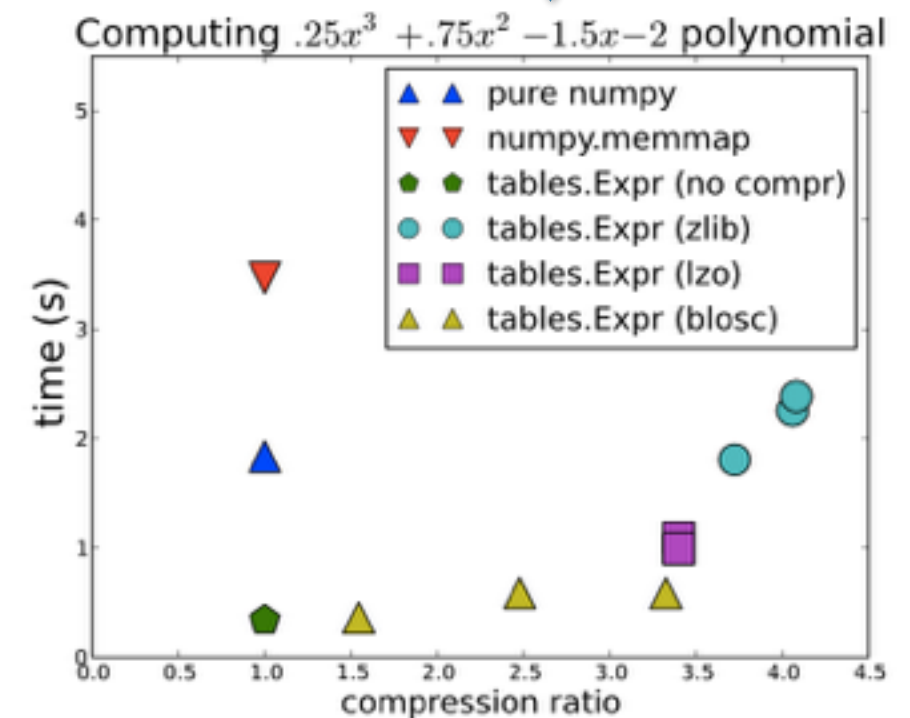
- Nice way to realize yourself while helping others




**The technology platform
to make a difference in
your relationship with
large and complex data**

HDF5 + a Twist

Out-of-core
Expressions



Indexed
Queries



Overview


- Recent trends in computer architecture
- The need for speed: storing and processing as much data as possible with your existing resources
- **Blosc & bcolz** as examples of compressor and data containers for large datasets that follow the principles of the newer computer architectures

Trends in Computer Storage

Forthcoming Trends

The growing gap between DRAM and HDD is facilitating the introduction of new SDD devices

The DRAM/HDD Speed Gap



From: *Solid State Drives in the Enterprise*
by Objective Analysis



PCIe SSD



M.2 SSD



BGA SSD

Latency Numbers Every Programmer Should Know


Latency Comparison Numbers

L1 cache reference	0.5	ns			
Branch mispredict	5	ns			
L2 cache reference	7	ns			14x L1 cache
Mutex lock/unlock	25	ns			
Main memory reference	100	ns			20x L2 cache, 200x L1 cache
Read 4K randomly from memory	1,000	ns	0.001	ms	
Compress 1K bytes with Zippy	3,000	ns			
Send 1K bytes over 1 Gbps network	10,000	ns	0.01	ms	
Read 4K randomly from SSD*	150,000	ns	0.15	ms	
Read 1 MB sequentially from memory	250,000	ns	0.25	ms	
Round trip within same datacenter	500,000	ns	0.5	ms	
Read 1 MB sequentially from SSD*	1,000,000	ns	1	ms	4X memory
Disk seek	10,000,000	ns	10	ms	20x datacenter roundtrip
Read 1 MB sequentially from disk	20,000,000	ns	20	ms	80x memory, 20X SSD
Send packet CA->Netherlands->CA	150,000,000	ns	150	ms	

Source: Jeff Dean and Peter Norvig (Google), with some additions

http://www.eecs.berkeley.edu/~rcs/research/interactive_latency.html

Reference Time vs Transmission Time



$t_{ref} \sim t_{trans} \Rightarrow$ optimizes storage access

Not All Storage Layers Are Created Equal

Memory: t_{ref} : 100 ns / t_{trans} (**1 KB**): ~100 ns

Solid State Disk: t_{ref} : 10 μ s / t_{trans} (**4 KB**): ~10 μ s

Mechanical Disk: t_{ref} : 10 ms / t_{trans} (**1 MB**): ~10 ms

The slower the media, the larger the block that is worth to transmit

But essentially, a blocked data access is mandatory for speed!

We Need More Data Blocking In Our Infrastructure!

- Not many data containers focused on blocking access
- No silver bullet: we won't be able to find a single container that makes everybody happy; it's all about tradeoffs
- With blocked access we can use persistent media (disk) as it is ephemeral (memory) and the other way around -> independency of media!

Can We Get Better
Bandwidth Than
Hardware Allows?

Compression for Random & Sequential Access in SSDs

Performance Specification	Incompressible Data	Compressible Data
Sequential Write Bandwidth (Mbps)	235	520
Sequential Read Bandwidth (Mbps)	550	550
Random Write (IOPS)	16,500 (65MB/s)	60,000 (240MB/s)
Random Read (IOPS)	46,000 (180MB/s)	50,000 (200MB/s)

3. Source: Intel® Solid-State Drive 520 Series Product Specification; Random reads based on 4KB Queue Depth 32

- Compression does help performance!

Compression for Random & Sequential Access in SSDs

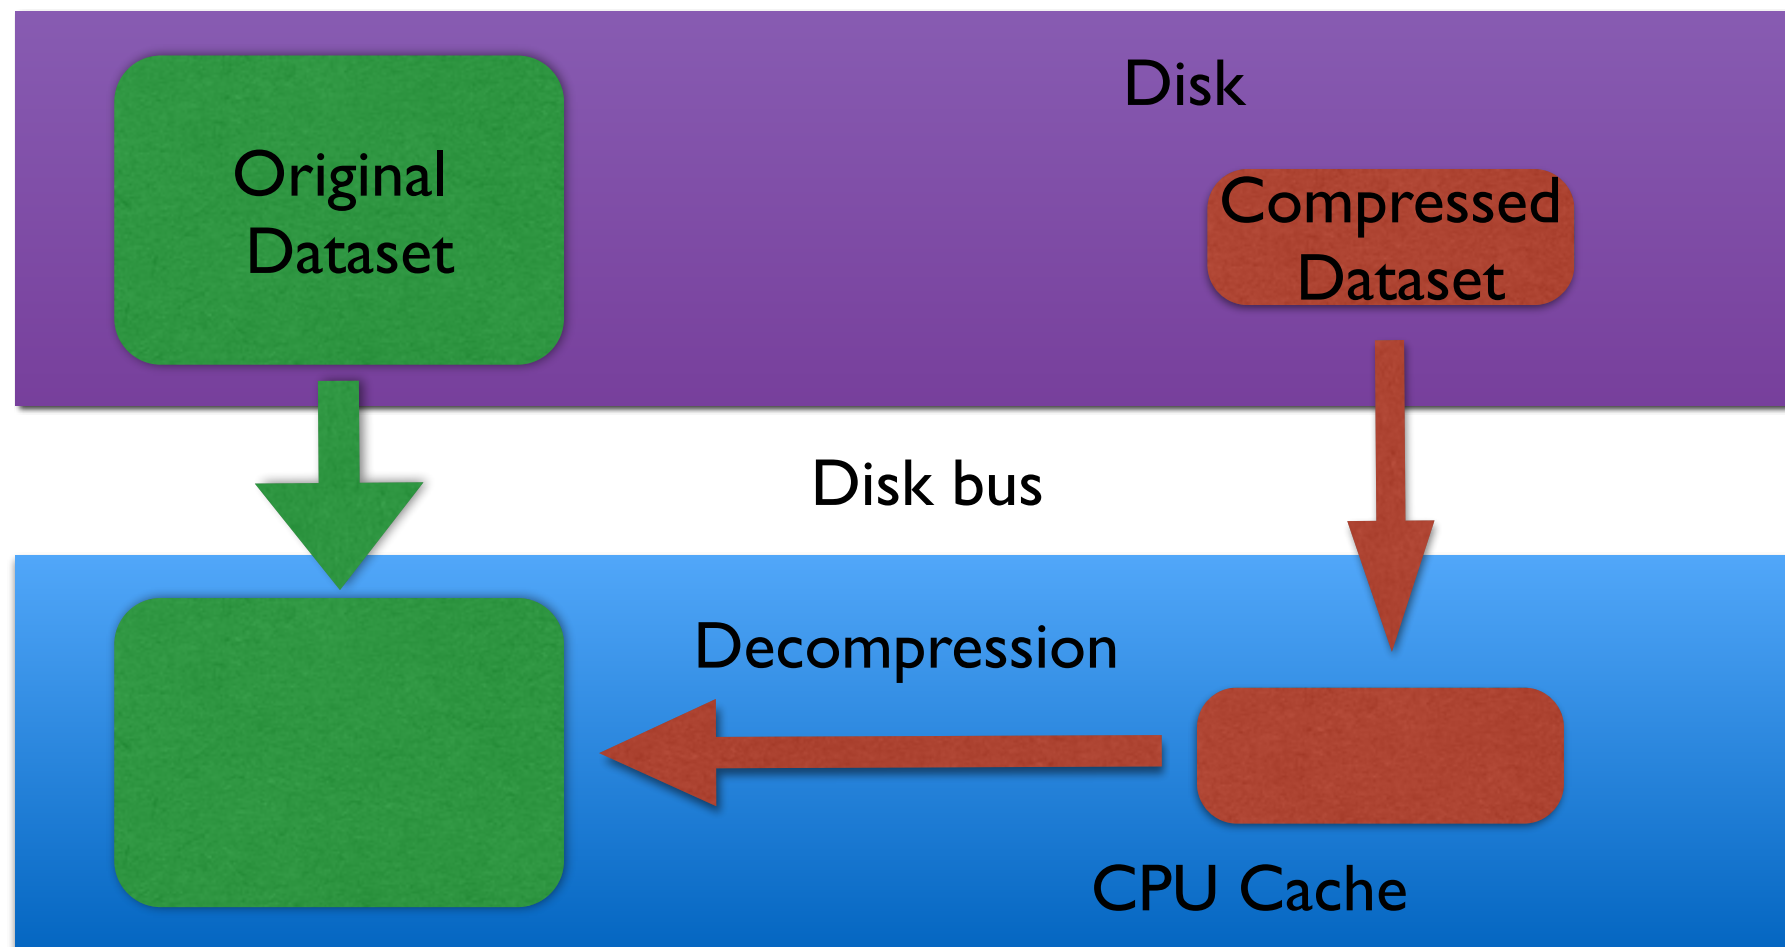
Performance Specification	Incompressible Data	Compressible Data
Sequential Write Bandwidth (Mbps)	235	520
Sequential Read Bandwidth (Mbps)	550	550
Random Write (IOPS)	16,500 (65MB/s)	60,000 (240MB/s)
Random Read (IOPS)	46,000 (180MB/s)	50,000 (200MB/s)

3. Source: Intel® Solid-State Drive 520 Series Product Specification; Random reads based on 4KB Queue Depth 32


- Compression does help performance!
- However, limited by SATA bandwidth

Leveraging Compression Straight To CPU

Less data needs to be transmitted to the CPU



Transmission + decompression faster than direct transfer?




When we have a fast enough compressor we can get rid of the limitations of the bus bandwidth.

How to get maximum compression performance?

Recent Trends In Computer CPUs

Memory Access Time vs CPU Cycle Time



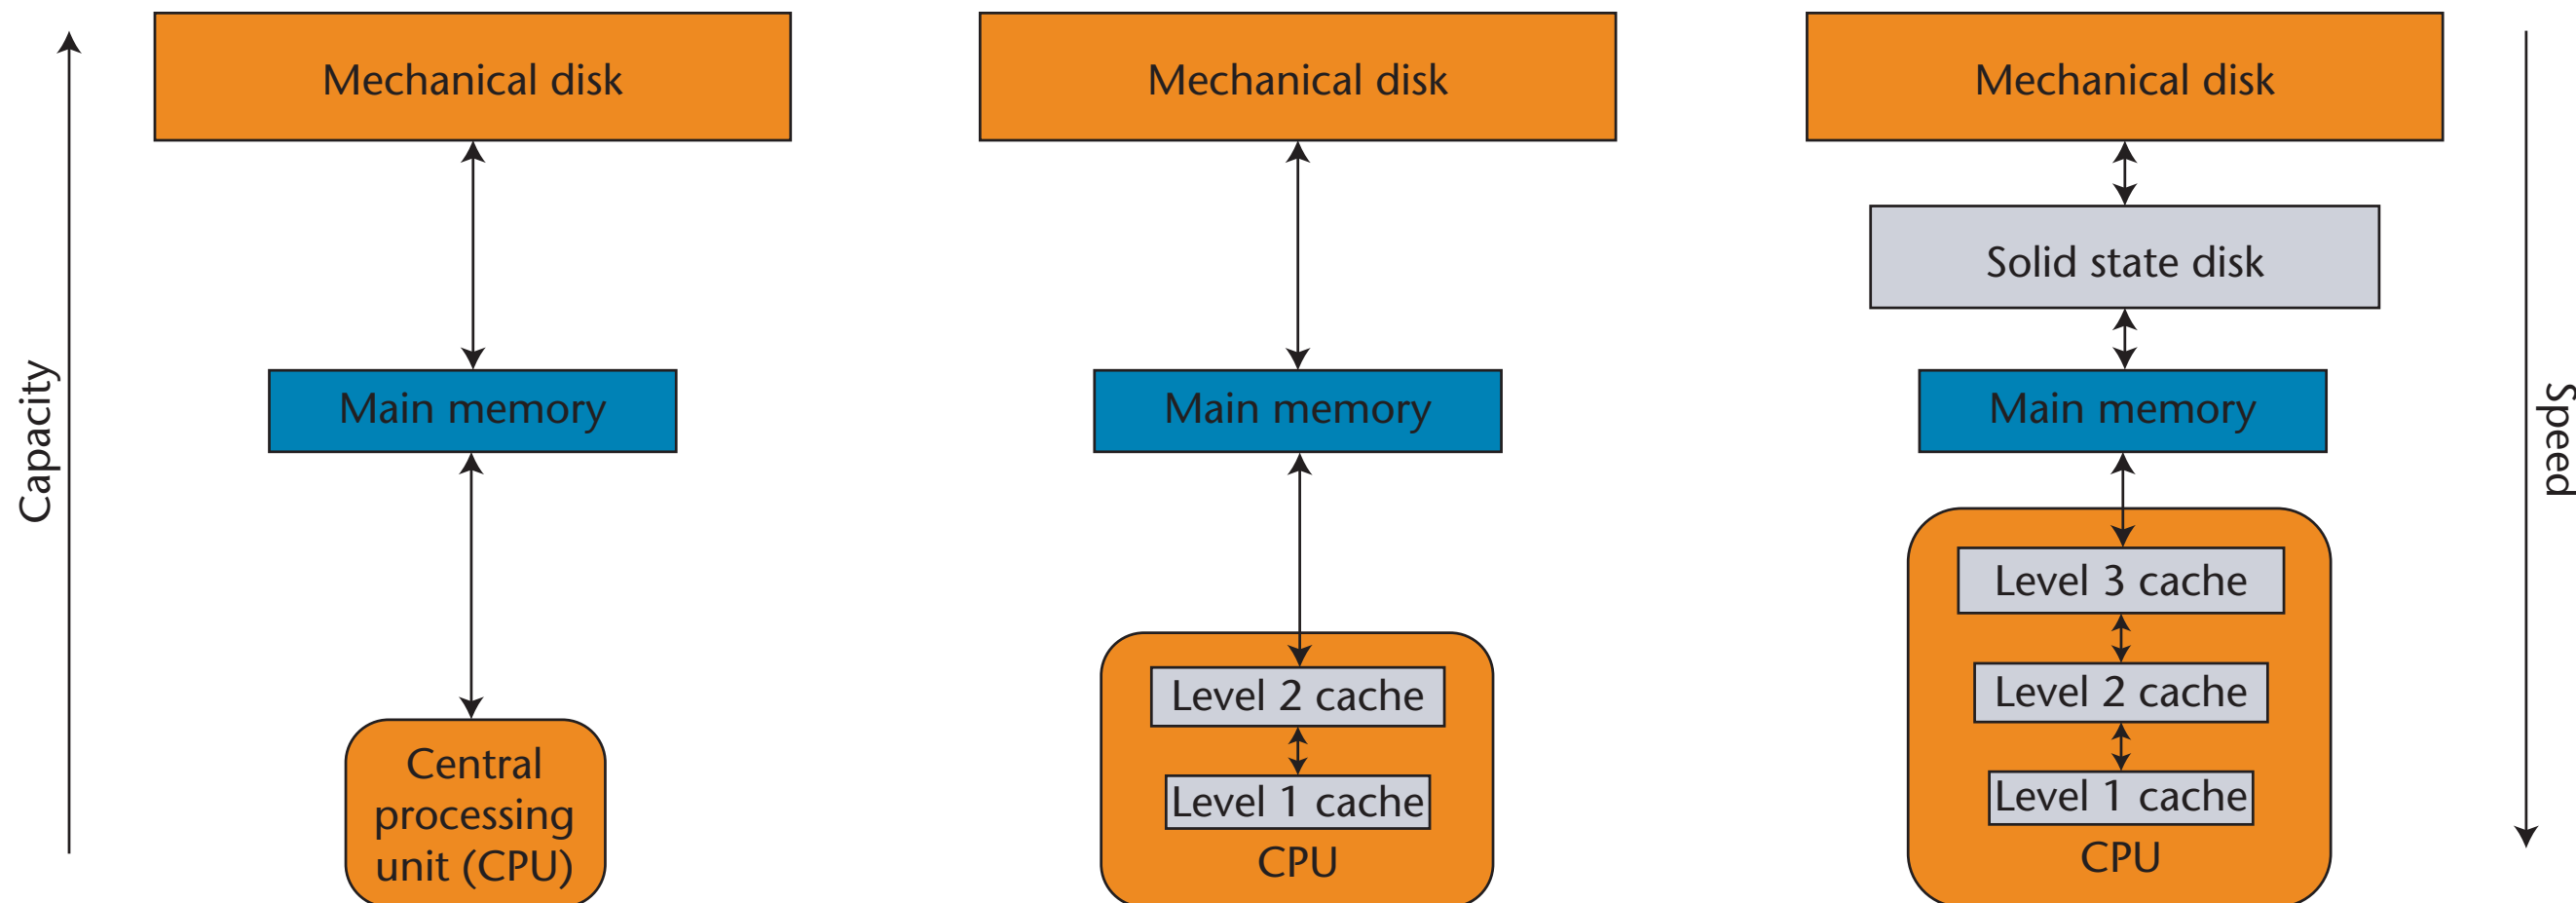
The gap is wide and still opening!

Computer Architecture Evolution

Up to end 80's


90's and 2000's

2010's




Hierarchy of Memory

By 2018 (Educated Guess)




Forthcoming Trends


More Cores




Wider Vectors



CPU+GPU Integration




Blosc: Compressing Faster Than *memcpy()*




Improving RAM Speed?

Less data needs to be transmitted to the CPU



Transmission + decompression faster than direct transfer?

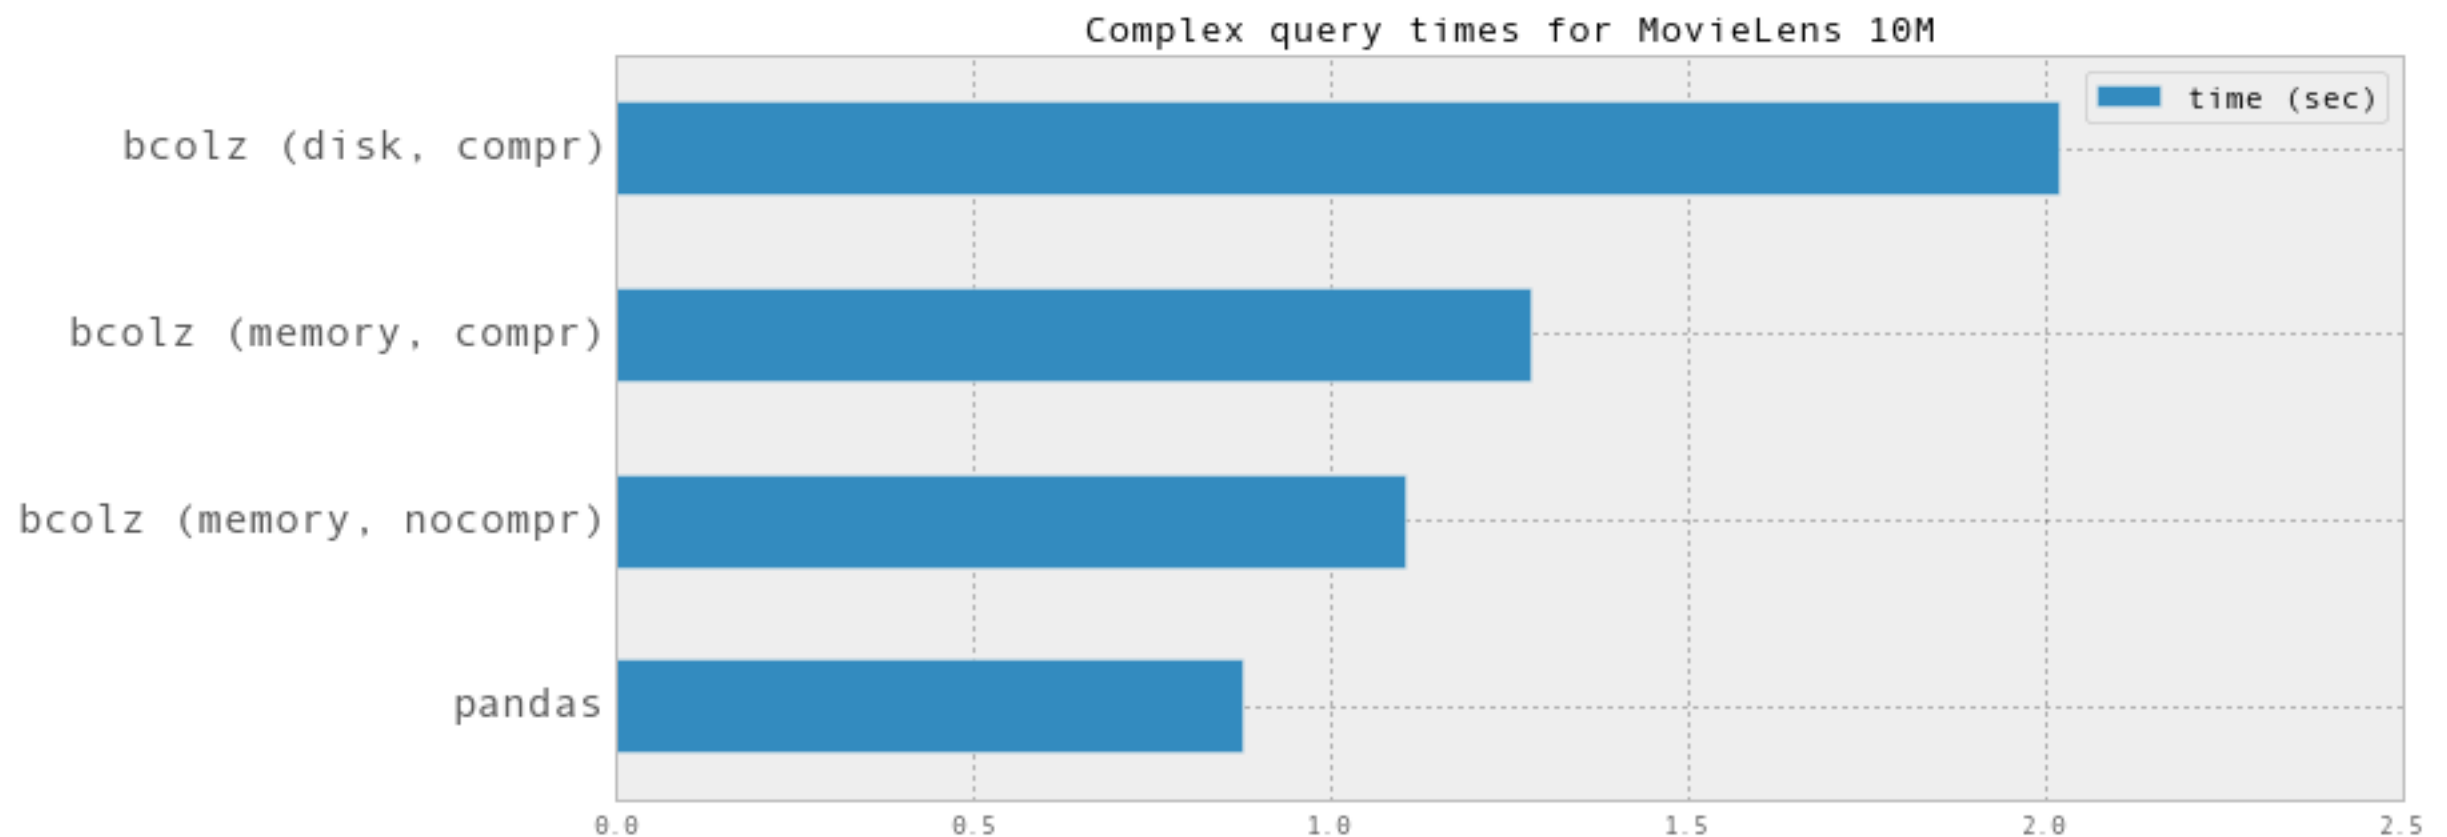
Source: <https://github.com/Blosc/movielens-bench>



Query Times


2012 old laptop (Intel Ivy-Bridge, 2 cores)
Compression **speeds** things up

Source: <https://github.com/Blosc/movielens-bench>



Query Times


2010 laptop (Intel Core2, 2 cores)
Compression still slow things down



bcolz vs pandas (size)

bcolz can store 20x more
data than pandas by using
compression

Accelerating I/O With Blosc



“Blosc compressors are the fastest ones out there at this point; there is no better publicly available option that I'm aware of. That's not just ‘yet another compressor library’ case.”

— *Ivan Smirnov*
(*advocating for Blosc inclusion in h5py*)

Compression matters!

Bcolz: An Example Of Data Containers Applying The Principles Of New Hardware

What is bcolz?

- bcolz provides data containers that can be used in a similar way than the ones in NumPy, Pandas
- The main difference is that data storage is **chunked**, not **contiguous**
- Two flavors:
 - **carray**: homogenous, n-dim data types
 - **ctable**: heterogeneous types, columnar


Contiguous vs Chunked

NumPy container



Contiguous memory

carray container

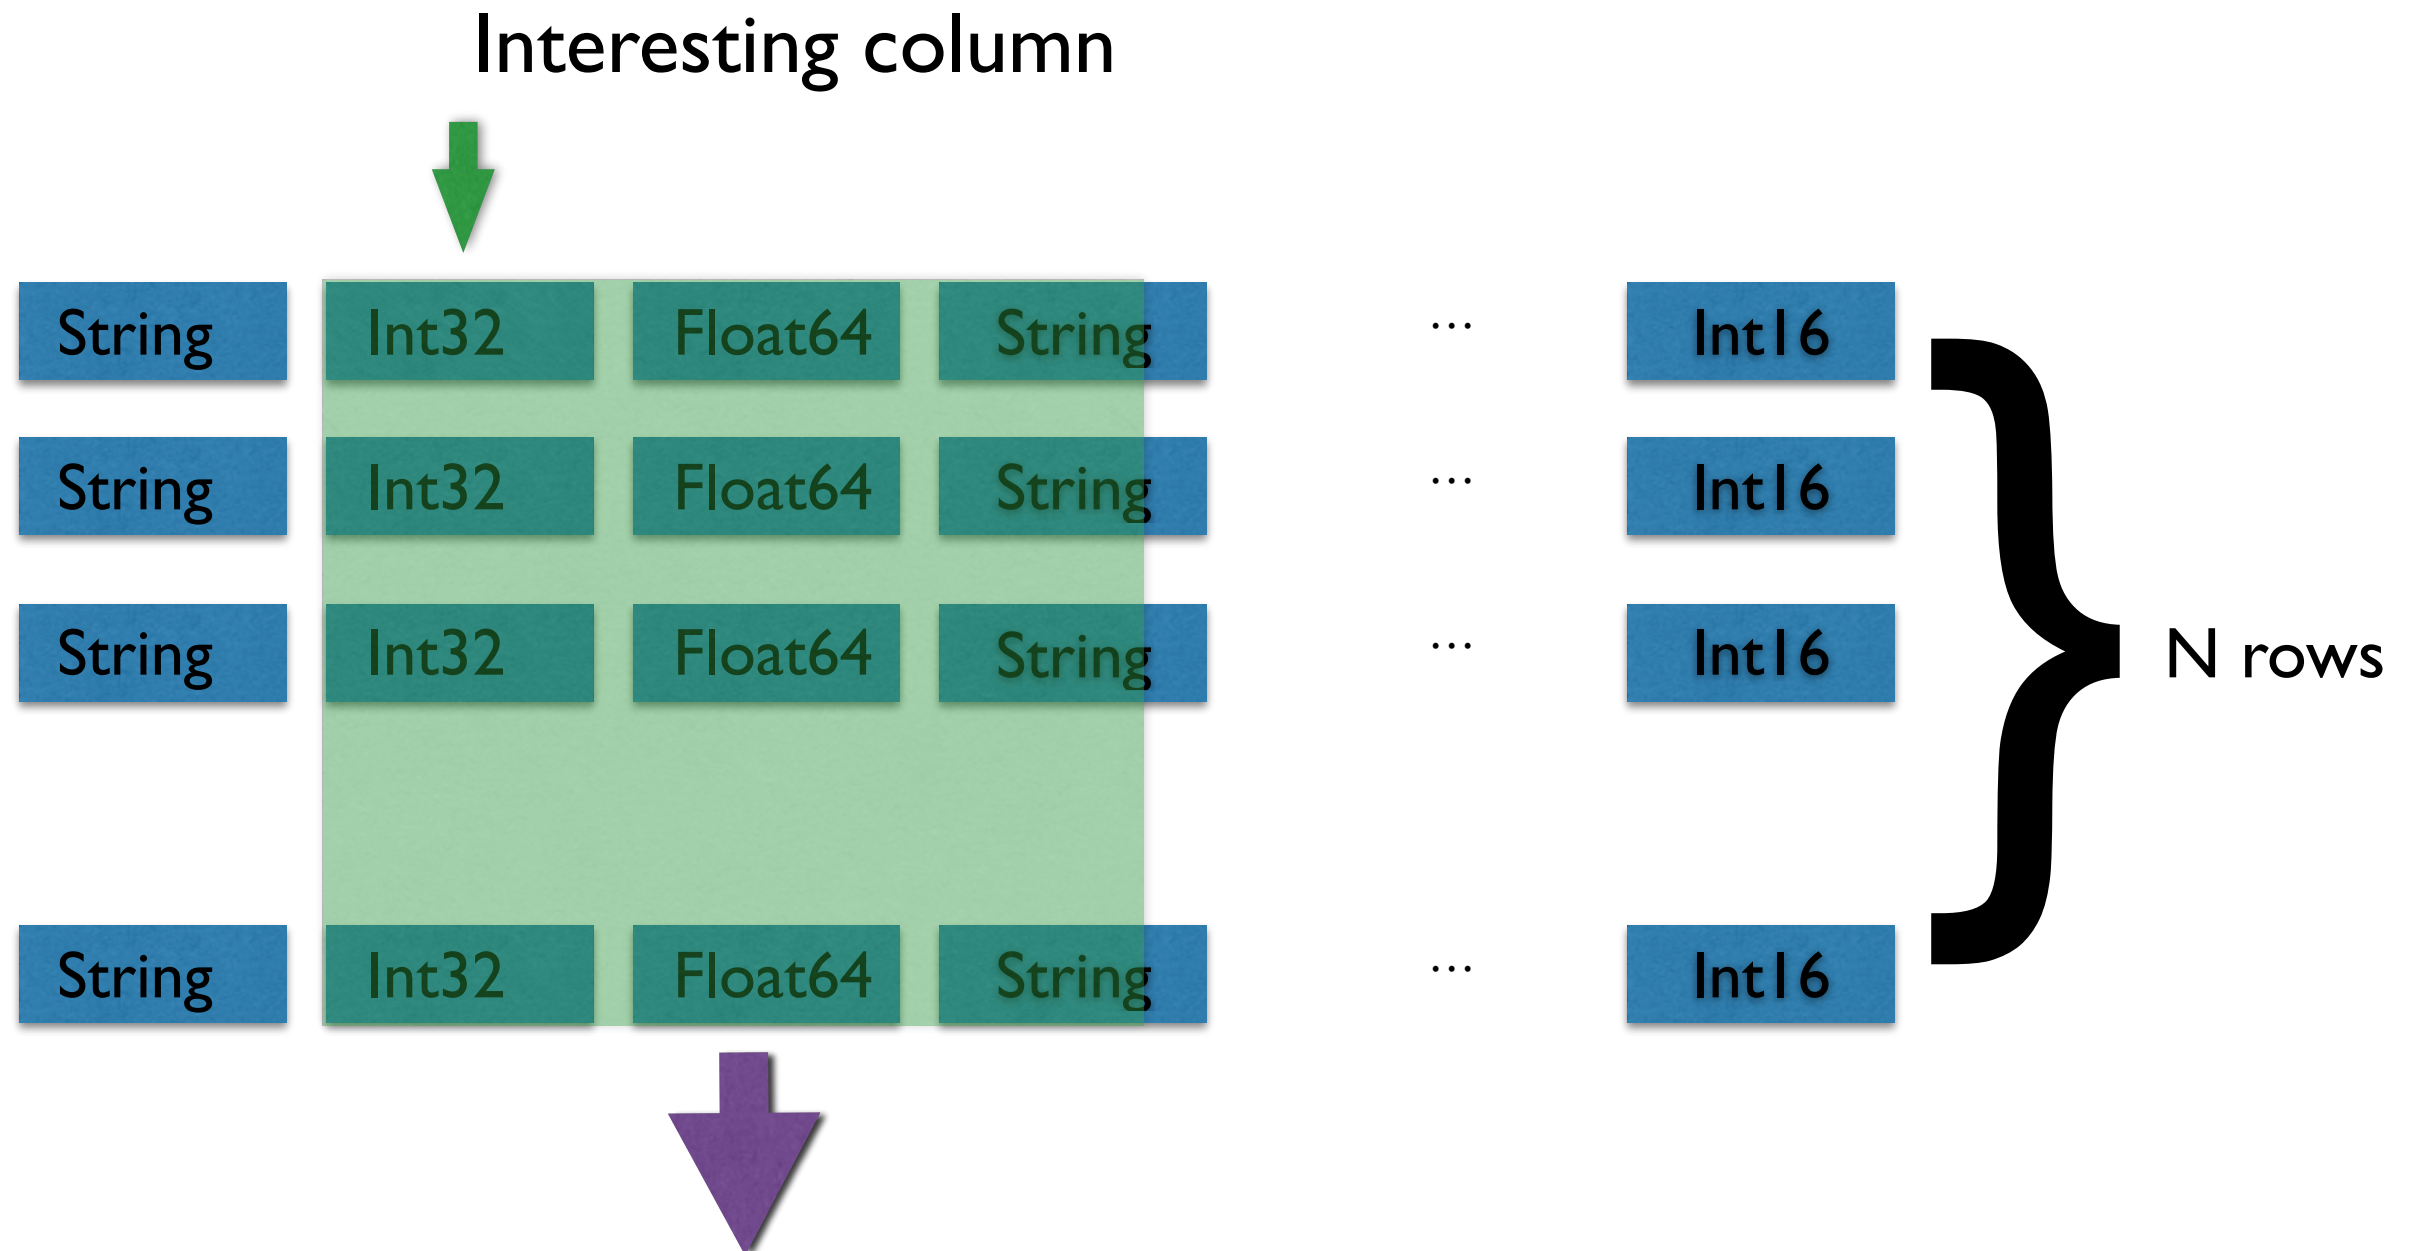


Discontiguous memory

Why Columnar?

- Because it adapts better to newer computer architectures

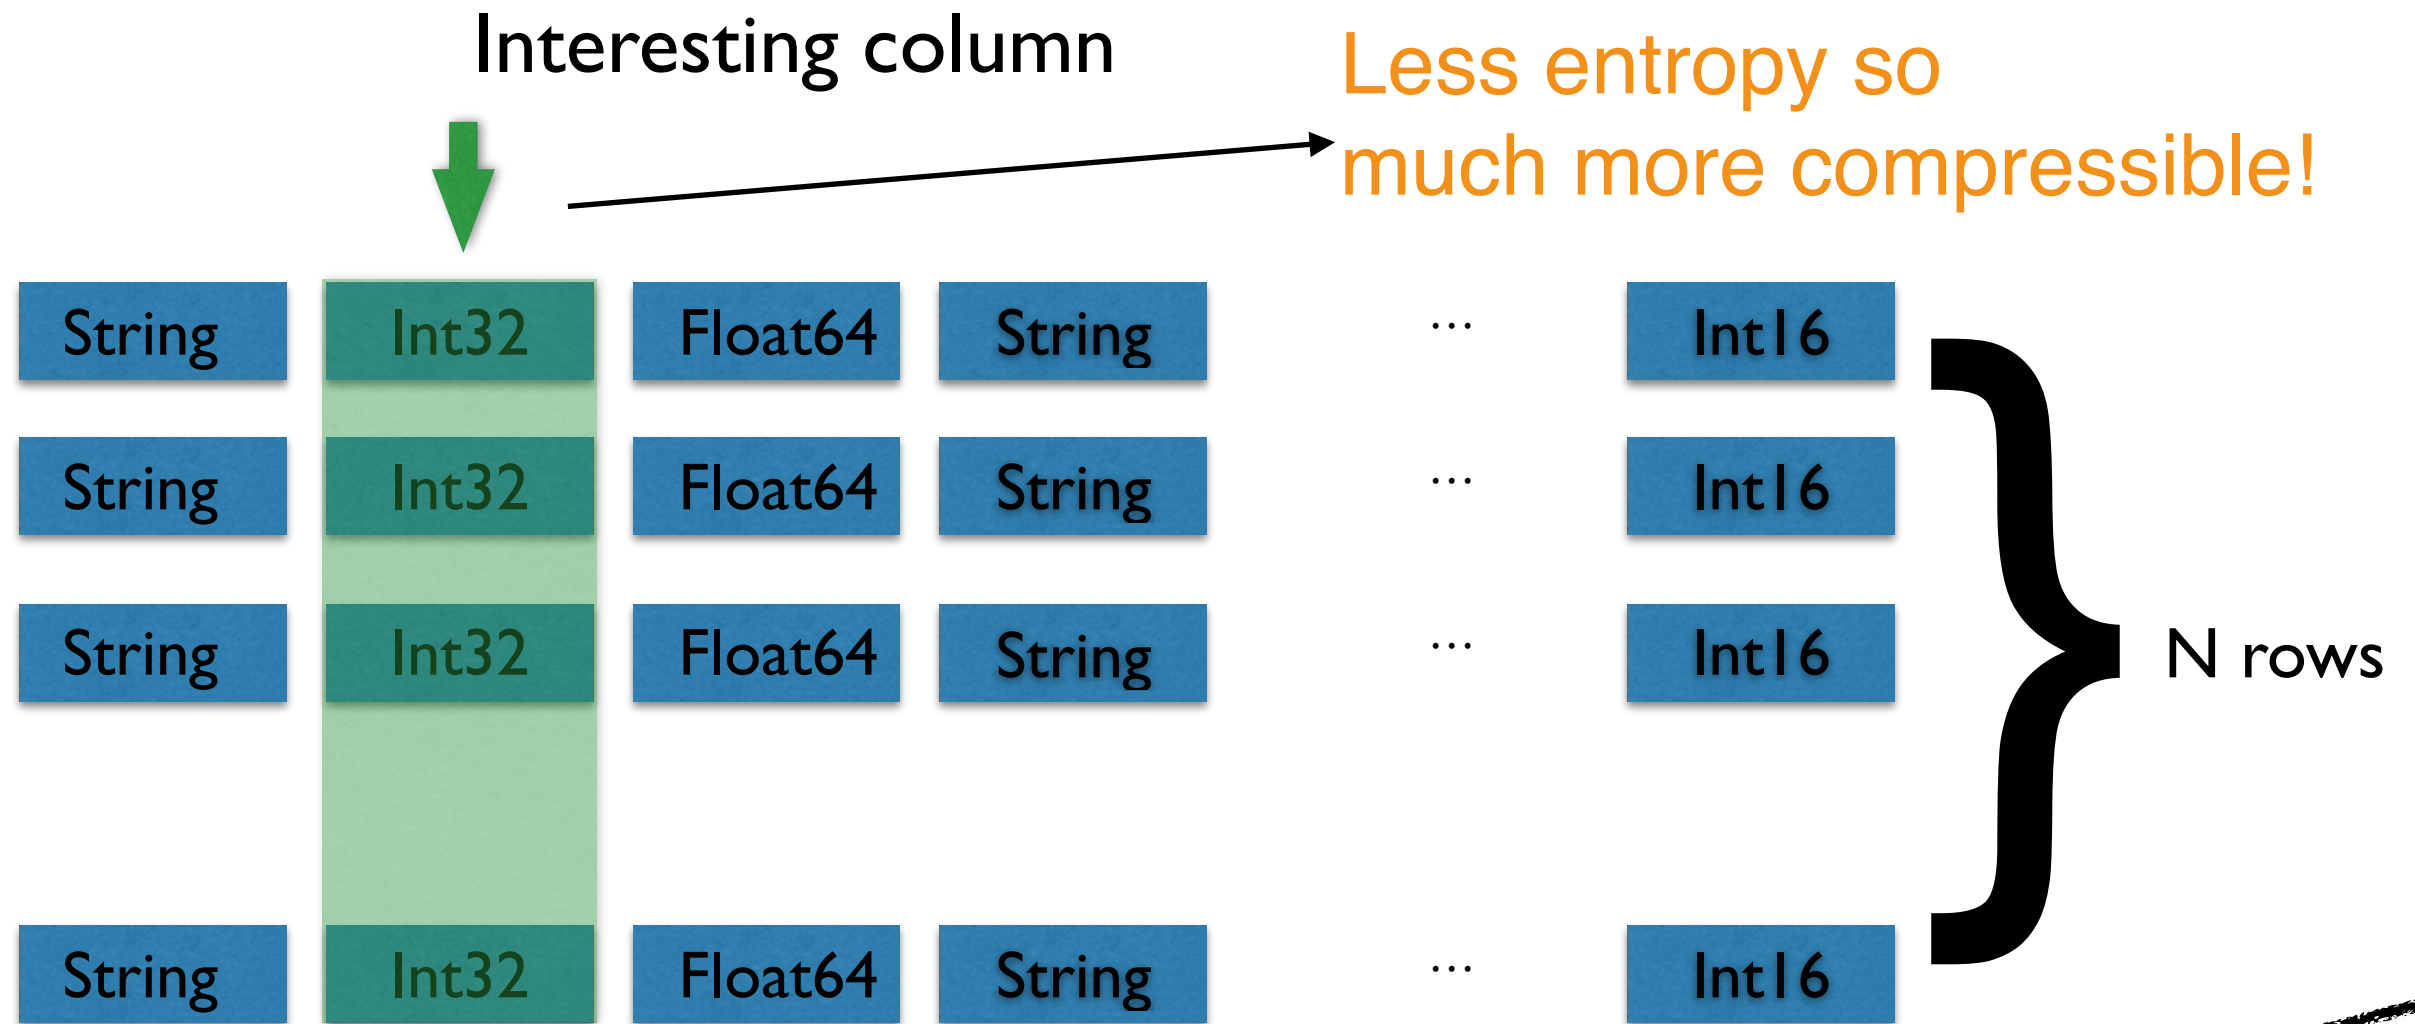
In-Memory Row-Wise Table (Structured NumPy array)



Interesting Data: $N * 4$ bytes (Int32)

Actual Data Read: $N * 64$ bytes (cache line)

In-Memory Column-Wise Table (bcolz *ctable*)




Interesting Data: $N * 4$ bytes (Int32)
Actual Data Read: $N * 4$ bytes (Int32)

Less memory travels
to CPU!

Some Projects Using bcolz

- Visualfabriq's bquery (out-of-core groupby's):
<https://github.com/visualfabriq/bquery>
- Scikit-allel:
<http://scikit-allel.readthedocs.org/>
- Quantopian:
<http://quantopian.github.io/talks/NeedForSpeed/slides#/>

bquery - On-Disk GroupBy

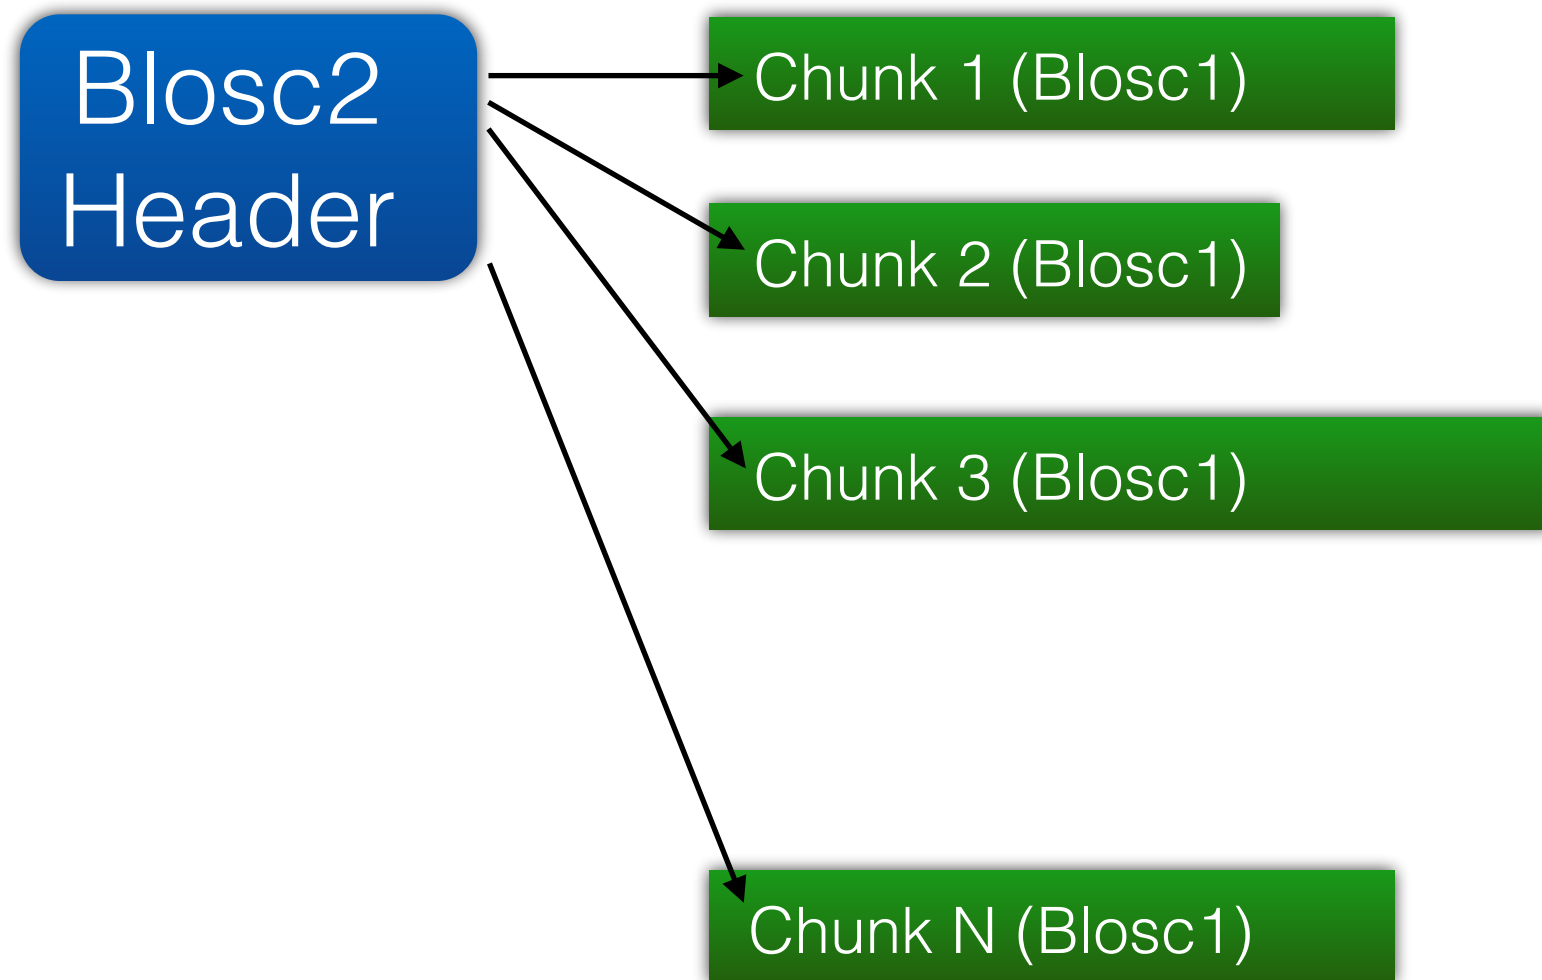


In-memory (pandas) vs on-disk (bquery+bcolz) groupby

“Switching to bcolz enabled us to have a much better scalable architecture yet with near in-memory performance”
— Carst Vaartjes, co-founder visualfabriq

“The future for me clearly involves lots of block-wise processing of multidimensional **bcolz** carrays”

–Alistair Miles
*Head of Epidemiological Informatics for the Kwiatkowski group. Author of **scikit-allel**.*





Introducing Blosc2

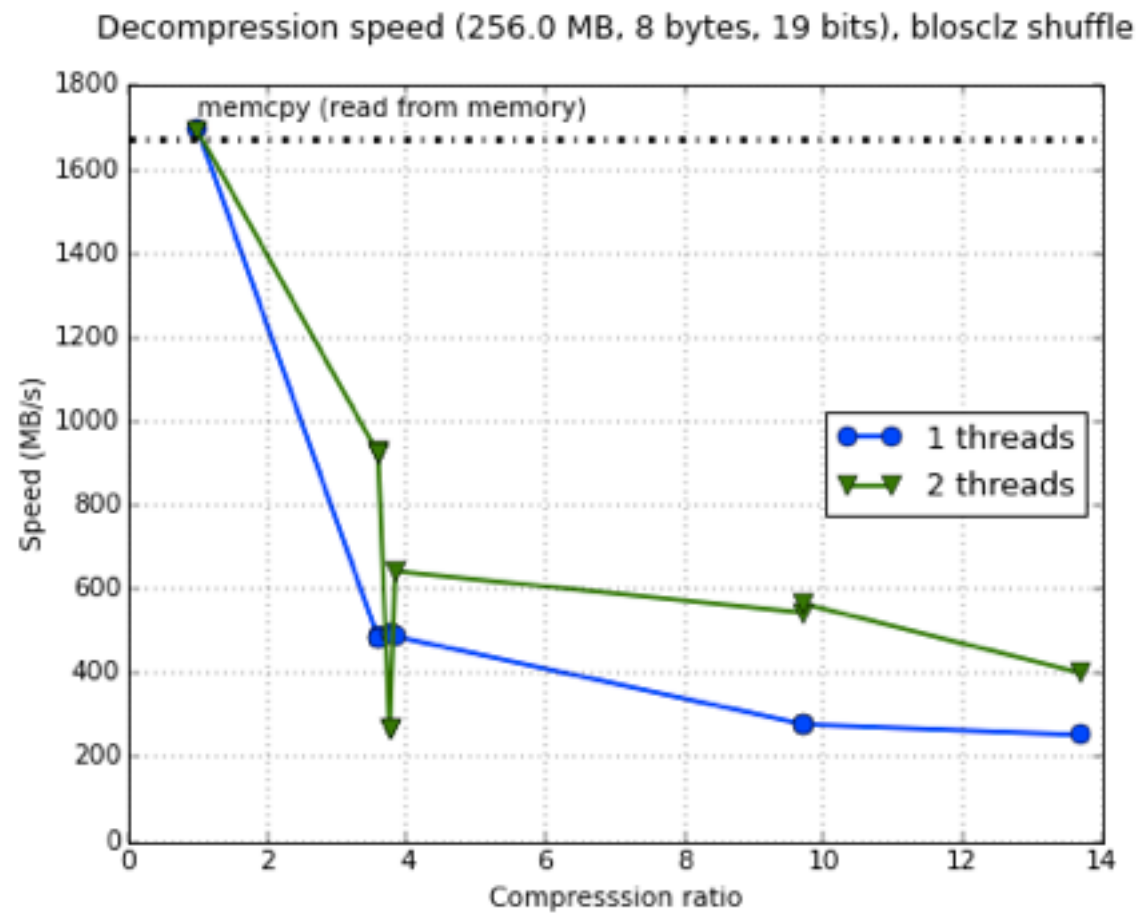
Next generation of Blosc

Planned features for Blosc2

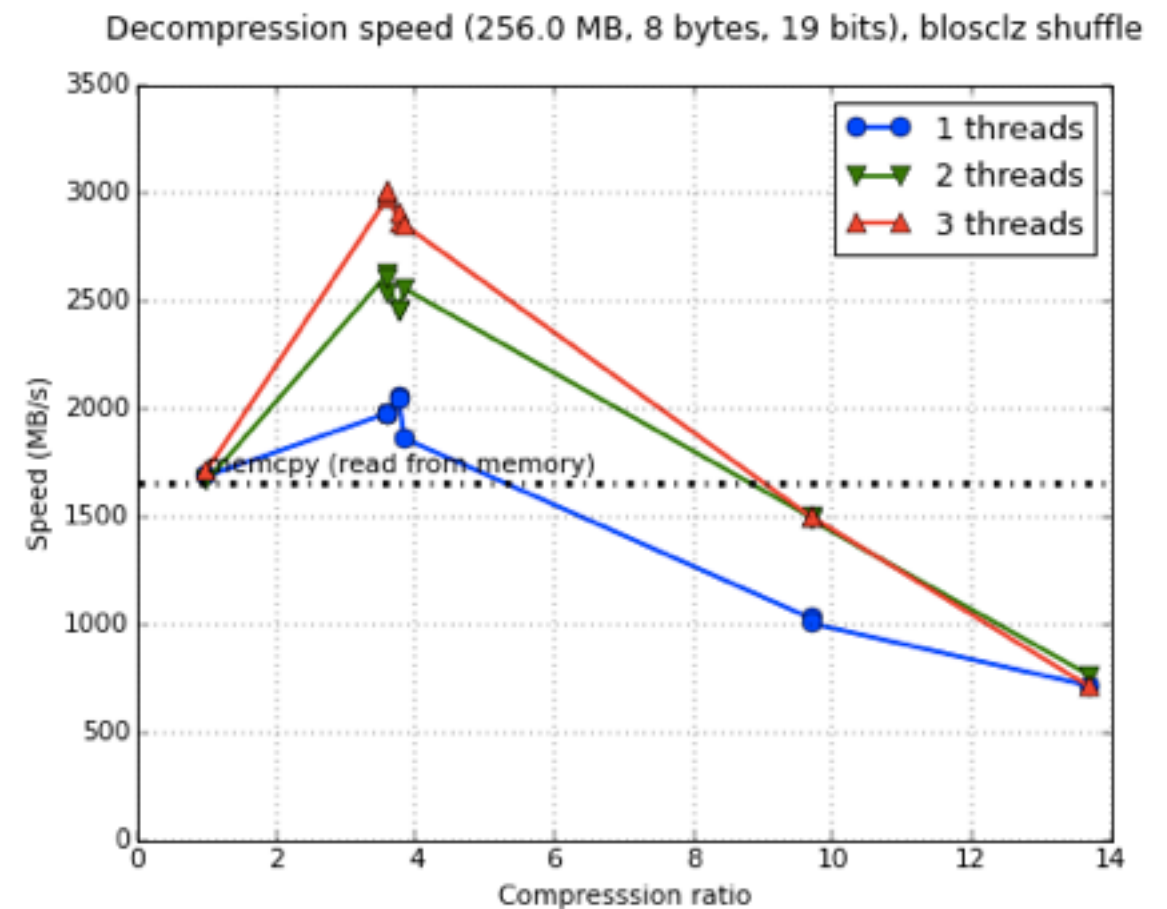
- Looking into inter-chunk redundancies (delta filter)
- Support for more codecs (Zstd is there already!)
- Serialized version of the super-chunk (disk, network)




Not using NEON



Using NEON



- At 3 GB/s, Blosc2 on ARM achieves one of the best bandwidth/Watt ratios in the market
- Profound implications for the density of data storage devices (e.g. arrays of disks driven by ARM)



Blosc2 has its own repo

<https://github.com/Blosc/c-blosc2>

Meant to be usable only when heavily tested!

(bcolz2 will follow after Blosc2)

Closing Notes

- Due to the evolution in computer architecture, the compression can be effective for two reasons:
 - We can work with more data using the same resources.
 - We can reduce the overhead of compression to near zero, and even beyond than that!

“In science, one can learn the most by studying what seems the least.”

–Marvin Minsky

¡Gracias!